

Tutorial 10 Query Processing

CSCI3170 Tutorial

Table of Content

- Query Evaluation
 - Join
 - Selection
 - Projection
- Query Optimization

Schema

- Supplier(sid, sname, location)
 - 500 pages, 80 tuples/page
- Supplier_Part(sid, pid, quantity)
 - 1000 pages, 120 tuples/page

Example

Select *

From Supplier S, Supplier_Part SP

Where **S.sid = SP.sid**

Join Operation

- Nested Loops Join
 - A tuple at a time
 - A page at a time
- Block Nested Loops Join
- Index Nested Loops Join
- Sort-Merge Join

A tuple at a time

```
for each tuples s in S do (S is called outer relation)  
  for each tuple sp in SP do (SP is called inner relation)  
    if (s.sid = sp.sid) then  
      add <s, sp> to result set
```

Cost:

- Scan S: 500 I/Os
- For each **tuple** of S, SP is scanned once: 1000 I/Os
- Total = 500 + **500 * 80 * 1000 = 40,000,500**
- Switch S and SP, the total is 1000 + 1000 * 120 * 500 = **60,001,000**

A page at a time

Improve the join by joining a page of tuples at a time

```
for each page p of S
  for each page q of SP
    output all s ∈ p and sp ∈ q such that s.sid = sp.sid
```


Cost

- Scan S: 500 I/Os
- For each **page** of S, SP is scanned once: 1000 I/Os
- Total: 500 + 500x1000 = **500,500**

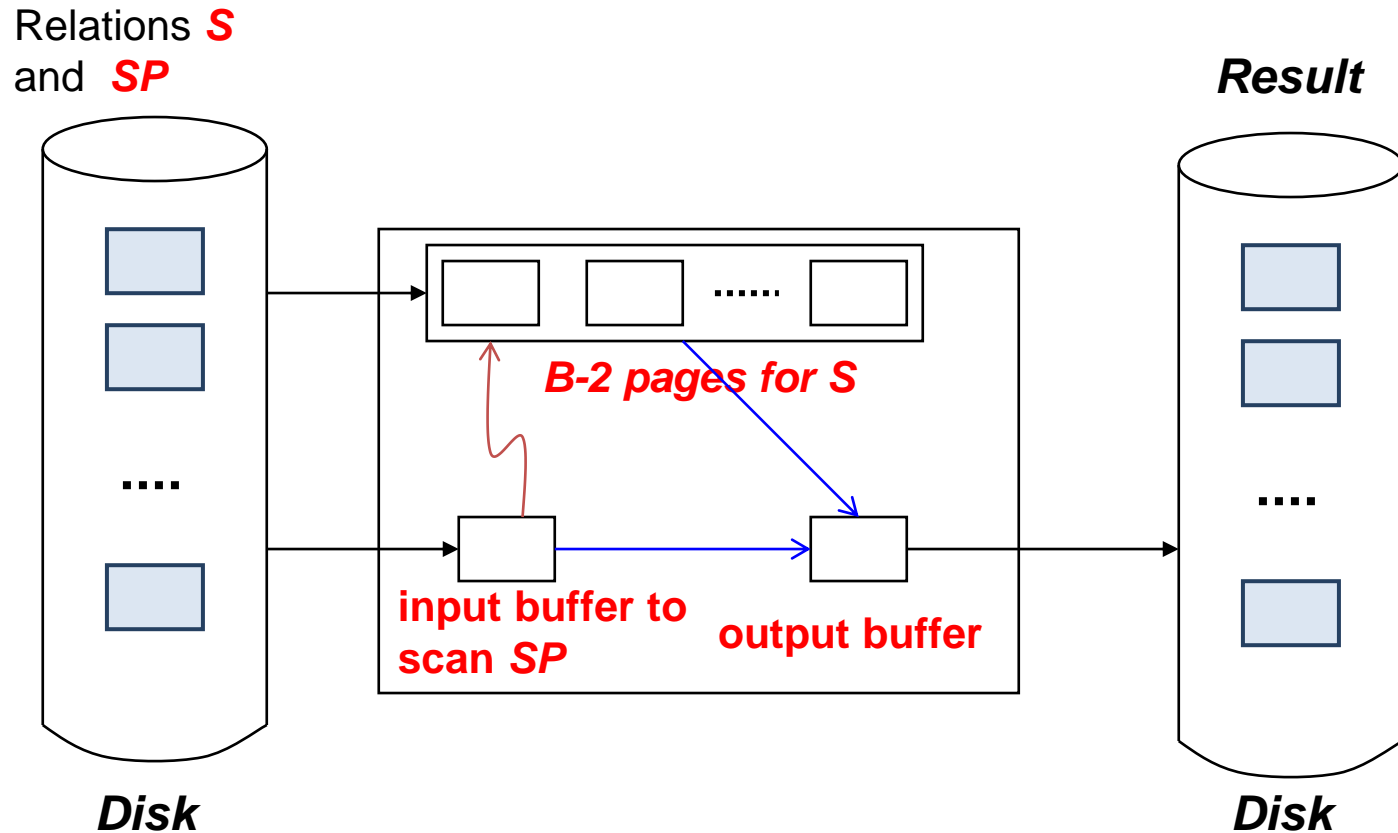
Block Nested Loops Join

For each block P for S
 For each page q for SP
 For each $s \in P$ and $sp \in q$ such that $s.sid = sp.sid$
 add $\langle s, sp \rangle$ to the result.

- Suppose we assign enough buffers to hold B pages. ($B > 2$)
 - 1 for input buffer (for SP)
 - 1 for output buffer
 - $B-2$ for outer relation (for S)



Block Nested Loops Join (Cont.)



Block Nested Loops Join (Cont.)

- Cost (Assume $B = 52$. So S is divided into 10 blocks.)
 - Scan S: 500 I/Os
 - For each **block** of S, scan SP once: 1000 I/Os
 - Total: $500 + 10 \times 1000 = 10,500$
 - Switch S and SP, the cost is: $1000 + 20 \times 500 = 11,000$
- Observation:
 - Choice of outer and inner relation will affect the cost.
 - Choose the **smaller one** as the outer relation
 - The buffer size will affect the cost
 - The bigger is the buffer, the fewer is the I/O cost
 - Trade off between space and time

Index Nested Loops Join

- Assume we have a hash index on sid of S then

For each $sp \in SP$ do
 For each $s \in S$ where $s.sid = sp.sid$ (use index)
 add $\langle s, sp \rangle$ to the result


- Note that for each tuple in S , we use index to find match tuple in SP .

Cost

Obtained from experiments
(for overflow pages)

- Scan SP : 1000 I/Os
- For each **tuple** in SP , an average of **1.2 I/O** to get to the bucket page containing the matching S data entry, retrieve the S tuple for 1 I/O (Note: sid is the primary key of Supplier relation)
- Total: $1000 + 120 \times 1000 \times (1+1.2) = 265,000$ I/Os

Index Nested Loops Join (2)



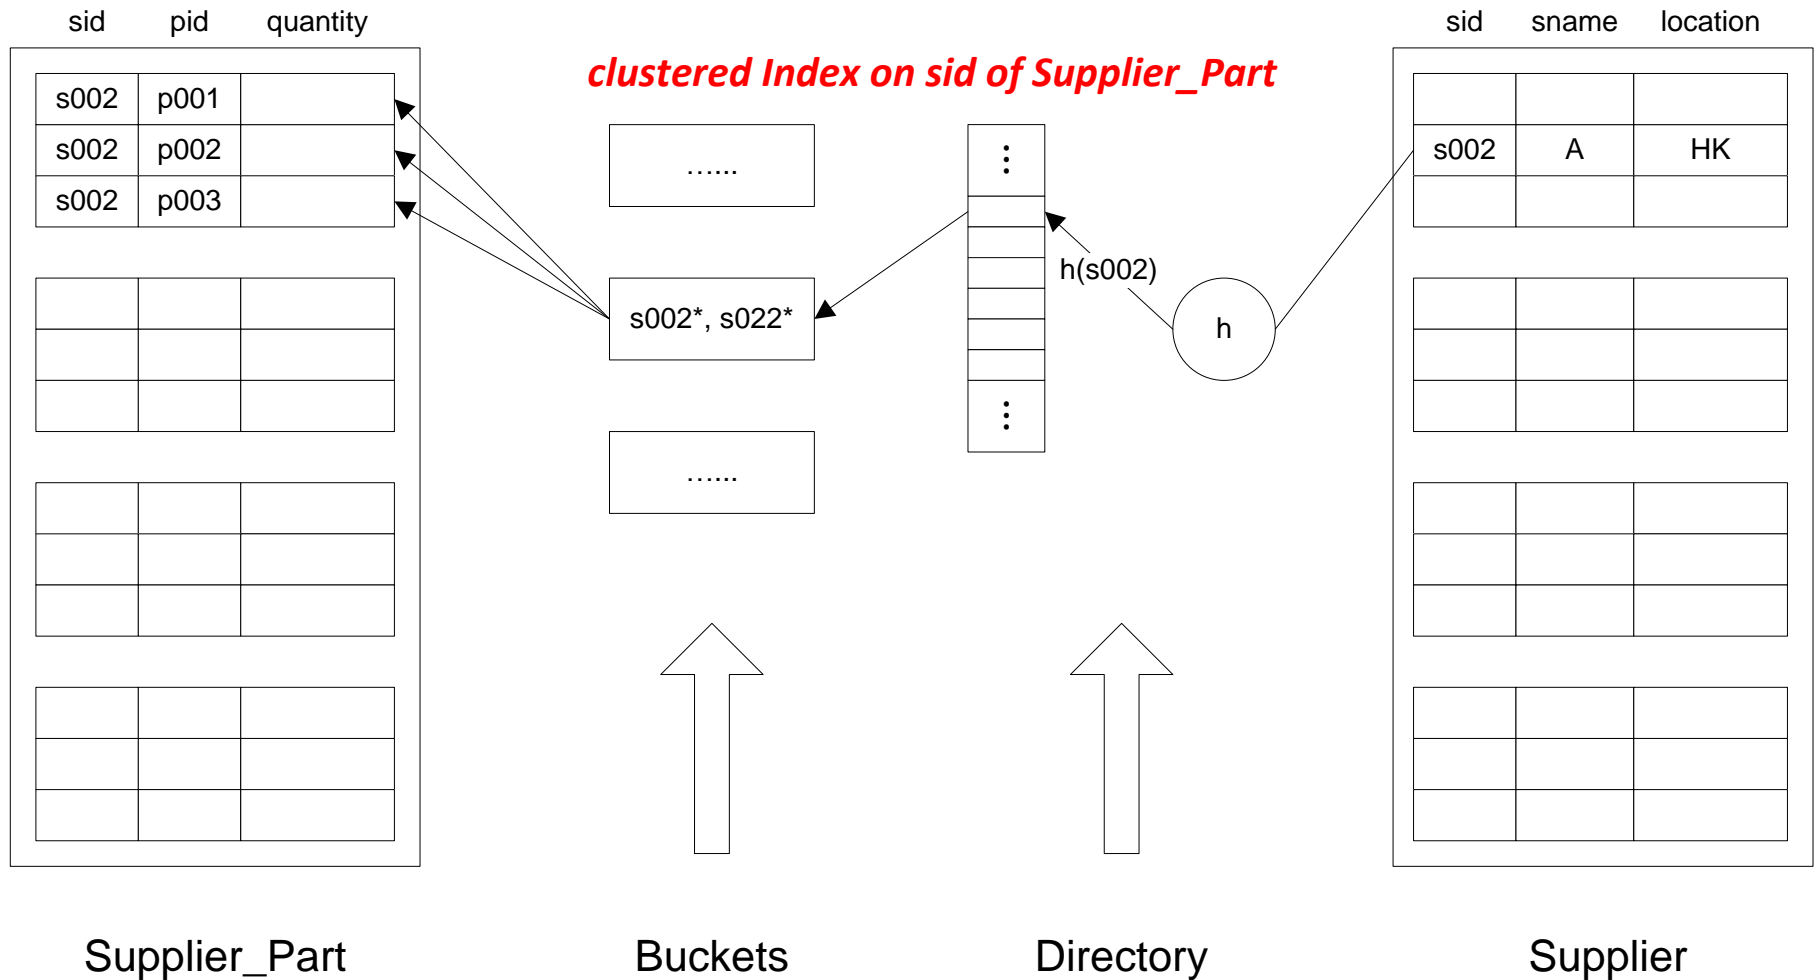
Index Nested Loops Join (3)

- Assume a hash index on sid of SP

```
foreach s ∈ S do  
    foreach sp ∈ SP where s.sid = sp.sid (use index)  
        add <s, sp> to the result
```


- Note that for each tuple in S, we use index to find match tuple in SP.

Index Nested Loops Join (4)



Index Nested Loops Join (5)

- Scan S: 500 I/Os
- For each **tuple** in S, an average of 1.2 I/O to get to the bucket page containing the matching SP data entry
- Estimation: 40000 suppliers supply 120000 parts, so each supplier supplies 3 parts on average.
 - Clustered Index (3 parts are in same page):
 $\text{total} = 500 + 40000 \times (1.2 + 1) = 88,500$
 - Unclustered Index (3 parts are in 3 different pages in the worst case):
 $\text{total} = 500 + 40000 \times (1.2 + 3) = 168,500$



Sort-Merge Join


- $S \bowtie_{i=j} SP$
- Sort Supplier and Supplier_Part in ascending order on the sid, then scan them to do a merge
- Scan S until s.sid \geq sp.sid
- Scan SP until sp.sid \geq s.sid
- Until s.sid = sp.sid. At this point, all S tuples with same value in S_i (current S group) and all SP tuples with same value in SP_j (current SP group) match. Output $\langle s, sp \rangle$ for all pairs of such tuples.
- Resume scanning S and SP

Sort-Merge Join (Cont.)

Case 1:

Both join attributes are not the primary key

R1 is scanned once, R2 group is scanned once per matching R1 tuple




Case 2:

One of the join attribute is the primary key

Both S, SP are scanned once

Cost of join is the sum of the size of S and SP, plus the cost of sorting these two relations



Query Optimization

- Motivation
 - Ideal: find the best plan
 - Practical: avoid the worst plan
- Optimization steps
- Query Evaluation Plan
- An example

Optimization Steps

- A query is essentially treated as a σ - Π - \bowtie algebra expression
- Optimizing such a relational algebra expression involves two basic steps:
 - Enumerate alternative plans for expression evaluation.
 - Estimate the cost of each plan and choose the plan with the lowest cost.

Query Evaluation Plan

- An extended algebra **tree** with annotations
- Each node indicates the relational operator and the implementation method for the relational operator.
- Each edge points to where the input comes from

Example

```
Select S.sname
From Supplier S, Supplier_Part SP
Where S.sid = SP.sid and
      SP.pid > 'p800' and
      S.location = 'HK'
```


- Supplier:
 - 500 pages, 80 tuples/page
 - 50 possible locations, uniformly distributed
- Supplier_Part:
 - 1000 pages, 120 tuples/page
 - Max part id is p1000, min part id is p1, uniformly distributed

Example (cont.)

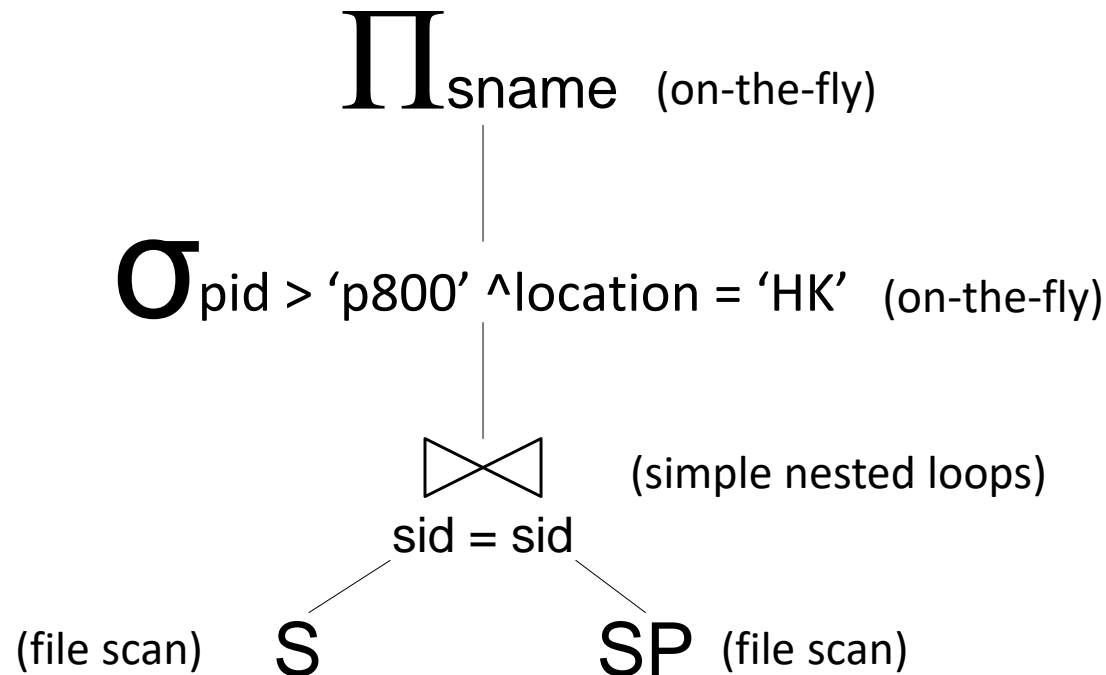
- Relation algebra of that query:

$$\Pi_{\text{sname}}(\sigma_{\text{pid} > \text{'p800'} \wedge \text{location} = \text{'HK'}}(S \bowtie_{\text{sid}=\text{sid}} SP))$$

- This algebra can be shown as a tree:



Full evaluation plan




Cost: $500 + 500 \times 1000 = 500500$

Push Selection and projection

- May greatly reduce the data involved in a join
- In the example, if we do the selection and projection first, then the join will involve small portion of data with sid, sname attributes only.
- May reduce the cost sometimes

Example



Cost:

Scan S: 500 I/Os; write T1: 10 I/Os.

Scan SP: 1000 I/Os; write T2: 200 I/Os.

Sort-merge join of T1 and T2: $10 + 200 + \text{I/Os for sorting T1 and T2}$.