

Tutorial 11

Coarse Review

CSCI3170 Tutorial

What's Important


- Why Database: Files vs. DBMS
- Database Modeling: Entity-relationship Model
- How it Works: Relational Algebra and SQL-based Application
- Make it Faster: Schema Refinement
- The Underlying: Storage vs. Memory, Indexes, and B+ Tree
- More: Hashing, Concurrency Control, and Recovery
- (Note: This short slide aims to help you organize what you've learned, please don't forget to review the more comprehensive lecture slide.)

Files vs. DBMS

- Suppose a company try to handle a large amount of data in files.
 - Application must have to handle large amount of data between main memory and secondary storage. Problems include: Buffering; Page-oriented access; Address handling
 - Special programs to answer different queries.
 - Must protect data from inconsistency due to multiple concurrent users.
 - Handle crash recovery.
 - Restore the system to a consistent state after a crash.
 - Security and access control.
- DBMS is a piece of software designed to make the above tasks easier.
- We can use the features of a DBMS to manage the data in a robust and efficient manner.

Entity-Relationship diagram (E-R diagram)

- The E-R model can be presented graphically by an E-R diagram.



Relational data model

- Most DBMS today are based on the relational data model.
- Relation
 - the central data description construct in this model.
 - It can be thought of as a set of records.
 - A table with rows and columns.
 - Row – a record
 - Column – field, attribute.

Relational data model

sid	name	login	age	gpa
53666	Jones	Jones@cs	18	3.4
53688	Smith	Smith@ee	18	3.2
53650	Smith	Smith@math	19	3.8
53831	Madayan	Madayan@music	11	1.8
53832	Guldu	guldu@music	12	2.0

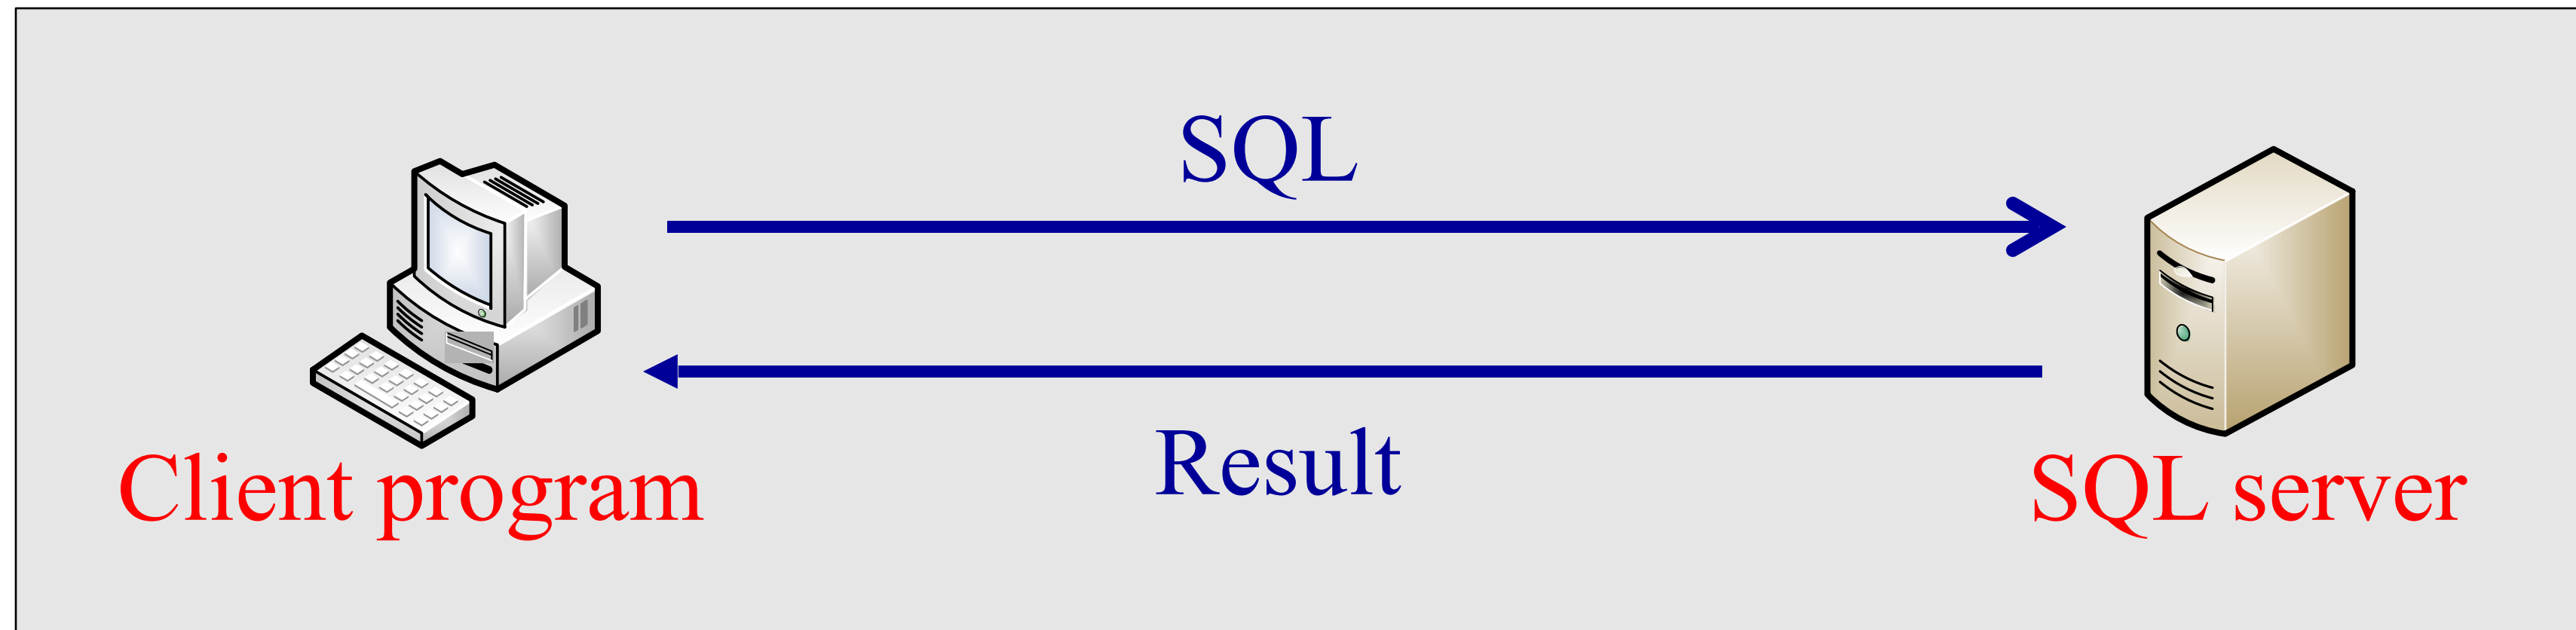
An example of a student relation

- A description of data in terms of a data model is called a schema.
- The schema of the above table is
Students(*sid*: string, *name*: string, *login*: string,
age: integer, *gpa*: real)

Basic Operators

- There are six basic operators in relation algebra:
 - select (σ)
 - project (Π)
 - union (\cup)
 - set different ($-$)
 - Cartesian product (\times)
 - rename (ρ)

Client-server Model




- The program running on the client machine sends SQL statements to the database server.
- The results of the SQL statements will be returned to the client.

Motivation Example

<u>Id</u>	name	age	rating	Hourly_wages	Hours_worked
123-22-3666	Peter	48	8	10	40
231-31-5368	Paul	22	8	10	30
131-24-3650	Mary	35	5	7	30
434-26-3751	David	35	5	7	32
612-67-4134	Ada	35	8	10	40

- Redundant Storage
 - The rating value 8 corresponds to the hourly wage 10, and this association is repeated three times.
 - Storage is not used efficiently.

<u>Id</u>	name	age	rating	Hourly_wages	Hours_worked
123-22-3666	Peter	48	8	10	40
231-31-5368	Paul	22	8	10	30
131-24-3650	Mary	35	5	7	30
434-26-3751	David	35	5	7	32
612-67-4134	Ada	35	8	10	40



<u>Id</u>	name	age	rating	Hours_worked
123-22-3666	Peter	48	8	40
231-31-5368	Paul	22	8	30
131-24-3650	Mary	35	5	30
434-26-3751	David	35	5	32
612-67-4134	Ada	35	8	40

rating	Hourly_wages
8	10
5	7

Functional dependency:
- rating determines Hourly_wages

A **decomposition of a relation schema** R consists of replacing the relation schema by two (or more) relation schemas that each contains a subset of attributes of R and together include all attributes in R

Disks and Files

- DBMS stores information on (“hard”) disks.
 - A disk is a sequence of bytes, each has a *disk address*.
 - **READ**: transfer data from disk to main memory (RAM).
 - **WRITE**: transfer data from RAM to disk.
 - Both are high-cost, relative to in-memory operations.
- Data are stored and retrieved in units called *disk blocks* or *pages*.
 - Each page has a fixed size, say 512 bytes. It contains a sequence of *records*.
 - In many (not always) cases, records in a page have the same size, say 100 bytes.
 - In many (not always) cases, records implement relational tuples.

Why Not Store Everything in Main Memory?

- *Costs too much.* Ram is much more expensive than disk.
- *Main memory is volatile.* We want data to be saved between runs. (Obviously!)
- Typical storage hierarchy:
 - Main memory (RAM) for currently used data.
 - Disk for the main database (secondary storage).
 - Tapes for archiving older versions of the data (tertiary storage).


Indexes

- An *index* on a file speeds up selections on the *search key fields* for the index.
 - Any subset of the fields of a relation can be the search key for an index on the relation.
 - *Search key* is *not* the same as *key* (minimal set of fields that uniquely identify a record in a relation).
- An index contains a collection of *data entries*
- A data entry is denoted as \mathbf{k}^* , where \mathbf{k} is a search key value and $*$ tells where to find the record containing \mathbf{k}
- Index must support efficient retrieval of all data entries \mathbf{k}^* with a given key value \mathbf{k}
- Structure of data entry in more detail

Index Classification

- *Primary vs. secondary*: If search key contains primary key, then called primary index, otherwise secondary.
 - *Unique* index: Search key contains a candidate key.
- *Clustered vs. unclustered*: If order of data records is the same as, or 'close to', order of data entries, then called clustered index.
 - A file can be clustered on at most one search key.
 - Cost of retrieving data records through index varies *greatly* based on whether index is clustered or not!

An example of a B+ tree

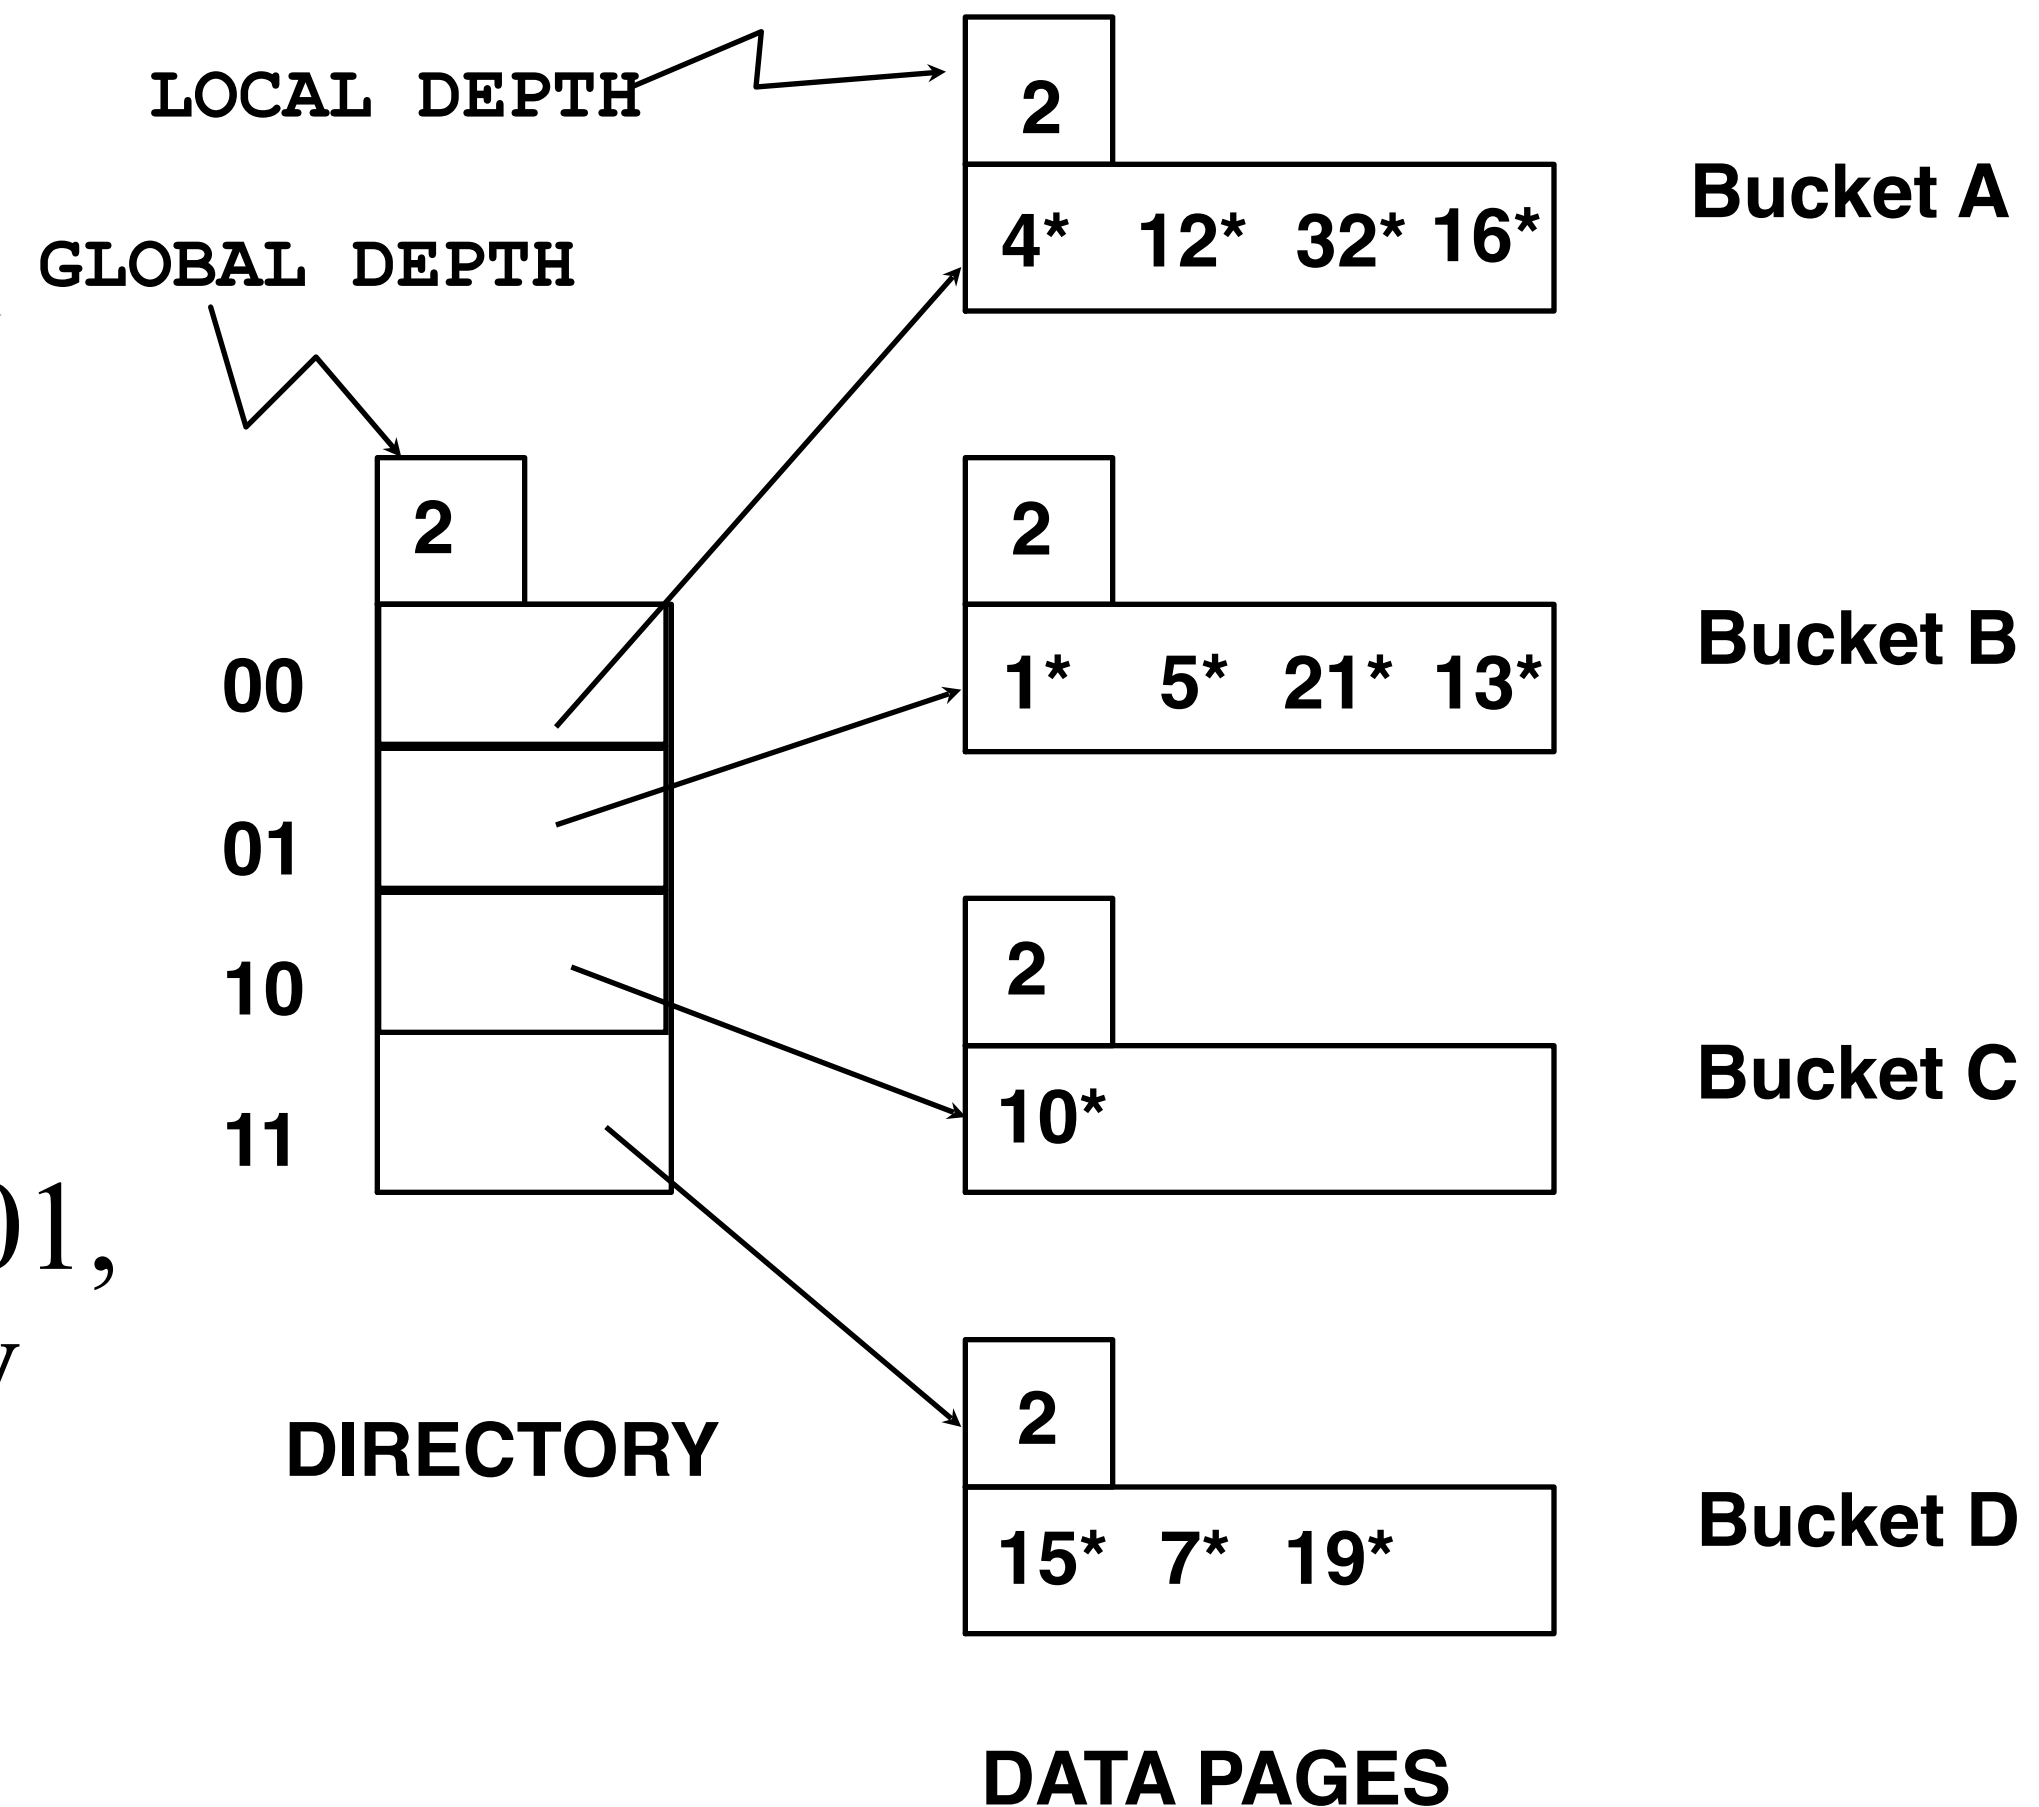


Extendible Hashing

- Situation: Bucket (primary page) becomes full. Why not re-organize file by *doubling* # of buckets?
 - Must re-hash all data entries to the right buckets
 - Example: assume hash function $h(k) = k \bmod M$
 - For $M = 4$, entries 3^* and 7^* both in bucket 3 ($3 \bmod 4 = 7 \bmod 4 = 3$)
 - But for $M = 8$, entry 7^* will be in bucket 7
 - Can we only re-hash those values that have changed addresses?
 - Difficulties: without re-hashing all the values, we don't know which values keep the old addresses and which get new addresses
 - Reading and writing all pages are expensive!
 - Question: how do we add more buckets, but only re-hash a few data entries?
 - Answer: use a level of indirection, *directory of pointers to buckets*

Example

- $h(r) = r \bmod 32$
- Directory is array of size 4.
- To find bucket for r , take last '*global depth*' # bits of $h(r)$
 - If $r = 5$, $h(r) = 5 = \text{binary } 101$, 5^* is in bucket pointed to by 01.



- ❖ **Insert:** If bucket is full, *split* it (allocate new page, re-distribute).
- ❖ *If necessary*, double the directory. (As we will see, splitting a bucket does not always require doubling; we can tell by comparing *global depth* with *local depth* for the split bucket.)

Concurrency

- Multiple users.
- Concurrent accesses.
- Problems could arise if there is no control.
- Example:
 - Transaction 1 (T1): withdraw \$500 from account A.
 - Transaction 2 (T2): deposit \$800 to account A.

```
Read(A)
A = A - 500
Write(A)
```

T_1

```
Read(A)
A = A + 800
Write(A)
```

T_2

Transactions

- A *transaction* is a sequence of read/write operations.
- A transaction is *atomic*.
 - Either all or none of the operations are executed.
 - No transaction can observe partial effect of other transactions.

Recoverable

- Transactions may be aborted due to logical failure. e.g. deadlock
- Recoverability is required to ensure that aborting a transaction does not change the semantics of committed transaction's operations.
- Example
 - Write₁(x,2); Read₂(x); Write₂(y,3); Commit₂
 - Not recoverable.
 - T₂ has committed before T₁ commits.
 - The problem is: what can we do if T₁ abort?
 - Delaying the commitment of T₂ can avoid this problem
- A schedule H is called *recoverable* (RC) if, whenever T_i reads from T_j. ($i \neq j$) in H and $c_i \in H$, $c_j < c_i$.
- Intuitively, a history is recoverable if each transaction commits after the commitment of all transactions (other than itself) from which it reads.